

Team 20 Code Sample doc

Environment setup

1. Spark installation: MAC:

<https://github.com/charlie-ph/BigDataAnalytics/blob/master/Installations-HowTos/How-To-Install-Spark-On-MACOS.md>

Windows:

<https://github.com/charlie-ph/BigDataAnalytics/blob/master/Installations-HowTos/How-To-Install-Spark-On-Windows.md>

2. PostgreSQL installation:

<https://www.postgresql.org/docs/current/tutorial-install.html>

How to run the code

Part 1: PostgreSQL

- 1) run code/Postgre/create_file.sql to create Data-warehouse in Postgre
- 2) import all csv files in /data into your database
- 3) run code/Postgre/SQL_basic_search.sql to get SQL results

Part2: Spark

- 1) run code/Spark/Example.ipynb in your notebook.
- 2) make sure you have correct address to access all csv files.

Dataset Explanation

I designed this database according to some basic requirements of my own trading system.

Req1 Data for all Instruments

Minute and day stock data, including open, close, high, low, etc.

Req 2 Company Basic Info

Basic information about the company including market capitalization, the company's industry, etc.

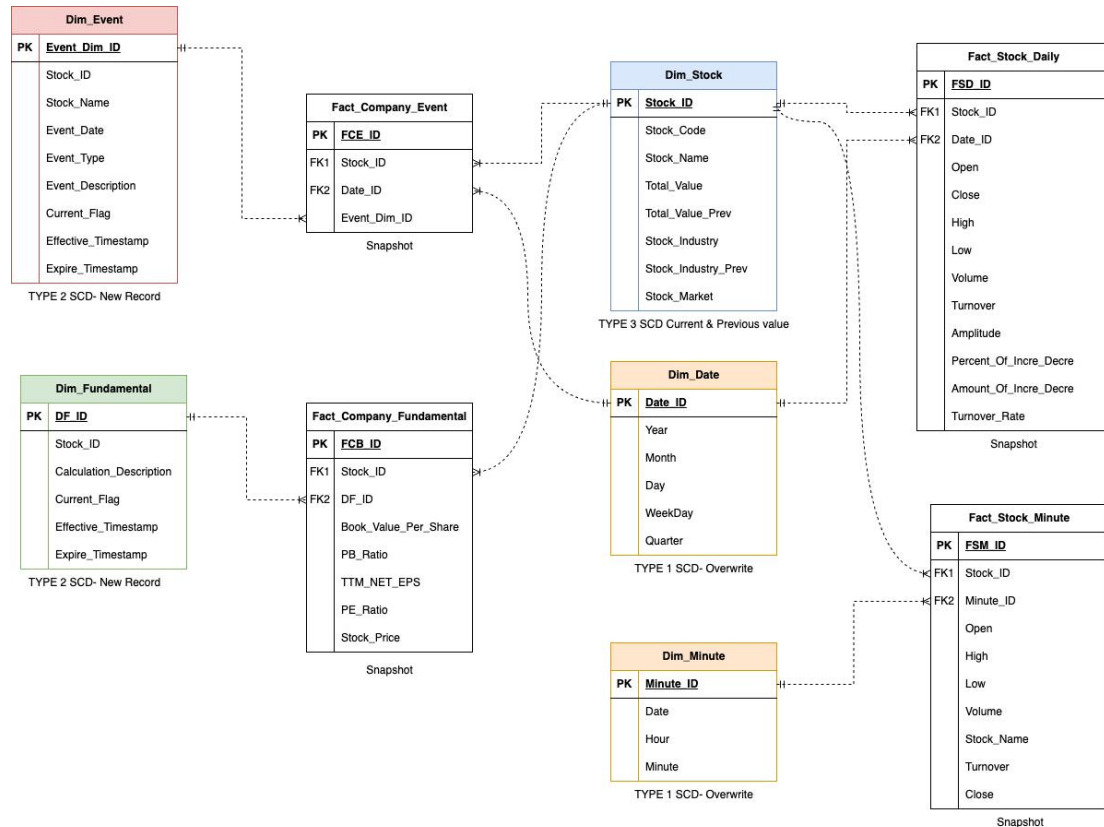
Req 3 Significant events occurring in the company(CN Market)

Significant events occurring in the company include asset reorganization, external guarantees, share pledges, asset acquisitions, etc.

Req 4 Company Earnings Data(US Market)

Earnings data released by the company each quarter, including PB, PE, etc.

Database ERD:



For each table, we provide one notebook to finish the ETL process(`code/Database/*.ipynb`). All the data is from this API: [akshare\(https://akshare.akfamily.xyz/\)](https://akshare.akfamily.xyz/)

Results of running the code with data

Seq1: select 10 stocks from Shanghai Stock Exchange(SSE)

PostgreSQL:

```

8  -- select any stock from Shanghai Stock Exchange(SSE)
9  SELECT ds.stock_id, ds.total_value from "Trading".Dim_Stock ds
10 where stock_id>600000
11 order by stock_id
12 limit 10;
13

```

Data Output Messages Explain × Notifications

	stock_id [PK] integer	total_value numeric (20,2)
1	600004	25820896467.53
2	600006	12820000000.00
3	600007	21102569087.30
4	600008	20113218454.98
5	600009	91352149991.40
6	600010	71739808751.84
7	600011	116165890856.60
8	600012	17150027400.00
9	600015	90715092267.60
10	600016	167686662862.66

Spark:

```

dim_stock_df = spark.read.csv('Dim_Stock.csv', header=True, inferSchema=True)
sse_stocks_df = dim_stock_df.filter(col("stock_id") > 600000).orderBy("stock_id").limit(10)
sse_stocks_df.select('Stock_ID', 'Total_Value').show()

```

0.3s

Stock_ID	Total_Value
600004	2.582089646753E10
600006	1.282E10
600007	2.11025690873E10
600008	2.011321845498000...
600009	9.135214999140001E10
600010	7.173980875184E10
600011	1.161658908566E11
600012	1.71500274E10
600015	9.07150922676E10
600016	1.676866628626E11

Seq2: select all stocks that rise over 9.9% on 2024-03-11

PostgreSQL:

```

14 --select all stocks that rise over 9.9% on 2024-03-11
15 SELECT * from "Trading".Fact_Stock_Daily
16 where Date_ID = 20240311 and percent_of_incre_decre>9.9
17 order by percent_of_incre_decre desc;
18
19 --select infomation from 20231120 9:30-9:40 on stock SH600004

```

	fsd_id [PK] bigint	stock_id integer	date_id integer	open numeric (10,2)	high numeric (10,2)	low numeric (10,2)	volume integer	turn nur
1	20240311300317	300317	20240311	4.04	4.84	4.04	972902	
2	20240311300324	300324	20240311	2.98	3.29	2.88	1842843	
3	20240311300530	300530	20240311	17.80	21.40	17.80	94640	
4	20240311300438	300438	20240311	22.01	25.91	22.01	425121	
5	20240311301202	301202	20240311	28.70	34.42	27.66	136343	
6	20240311301292	301292	20240311	14.16	16.85	14.16	50492	
7	20240311300890	300890	20240311	23.37	27.95	23.37	76936	
8	20240311301349	301349	20240311	28.89	34.66	28.89	30457	
9	20240311300712	300712	20240311	21.48	25.14	21.37	92134	
10	20240311688063	688063	20240311	85.00	99.60	85.00	136891	
11	20240311301205	301205	20240311	91.00	116.62	91.00	133827	

Spark:

```

fact_stock_daily_df = spark.read.csv('Fact_Stock_Daily_2024.csv', header=True, inferSchema=True)
using_stocks_df = fact_stock_daily_df.filter((col("Date_ID") == 20240311) & (col("percent_of_incre_decre") > 9.9))\
    .orderBy(col("percent_of_incre_decre").desc())
rising_stocks_df.show()

```

FSD_ID	Stock_ID	Date_ID	Open	High	Low	Volume	Turnover	Amplitude	Percent_Of_Incre_Dece	Amount_Of_Incre_Dece
20240311300317	300317	20240311	4.04	4.84	4.04	972902	4.60648926E8	19.85	20.1	0.81
20240311300324	300324	20240311	2.98	3.29	2.88	1842843	5.68963029E8	14.96	20.07	0.55
20240311300530	300530	20240311	17.8	21.4	17.8	94640	1.91813994E8	20.19	20.02	3.57
20240311301292	301292	20240311	14.16	16.85	14.16	50492	8.4292834E7	19.16	20.01	2.81
20240311301202	301202	20240311	28.7	34.42	27.66	136343	4.30519547E8	23.57	20.01	5.74
20240311300890	300890	20240311	23.37	27.95	23.37	76936	2.03484159E8	19.67	20.01	4.66
20240311301349	301349	20240311	28.89	34.66	28.89	30457	1.01086693E8	19.98	20.01	5.78
20240311300438	300438	20240311	22.01	25.91	22.01	425121	1.021499595E9	18.06	20.01	4.33
20240311688063	688063	20240311	85.0	99.6	85.0	136891	1.26584976E9	17.59	20.0	16.6
20240311300712	300712	20240311	21.48	25.14	21.37	92134	2.19739943E8	18.0	20.0	4.15
20240311301205	301205	20240311	91.0	116.62	91.0	133827	1.369576675E9	26.36	17.31	16.85
20240311300025	300025	20240311	8.54	10.08	8.38	306597	2.85149119E8	20.09	17.14	1.45
20240311300769	300769	20240311	41.01	46.89	41.01	302334	1.342825861E9	14.71	16.59	6.63
20240311300953	300953	20240311	45.8	53.0	45.8	31804	1.59406988E8	15.72	15.31	7.01
20240311688248	688248	20240311	27.78	31.31	27.78	171501	5.19394678E8	13.07	15.07	4.07
20240311300750	300750	20240311	165.0	181.65	164.1	763350	1.3248182793E10	11.11	14.46	22.85
20240311301236	301236	20240311	46.95	55.88	46.8	1141301	5.880365894E9	18.92	14.4	6.91
20240311300709	300709	20240311	38.8	44.98	36.89	446070	1.808201963E9	21.3	14.01	5.33
20240311300014	300014	20240311	38.6	42.76	38.55	648947	2.651262664E9	11.19	13.0	4.85
20240311300411	300411	20240311	8.12	9.49	8.12	962847	8.43633186E8	17.19	12.92	1.05

only showing top 20 rows

Seq3: select information from 20231120 9:30-9:40 on stock SH600004

PostgreSQL:

19	--select infomation from 20231120 9:30-9:40 on stock SH600004
20	SELECT * from "Trading".Fact_Stock_Minute
21	where Minute_ID>202311200930 and Minute_ID<202311200940 and stock_id = 600004;
22	
23	--show total minutes in 2024-03-11

Data Output	Messages	Explain ×	Notifications
-------------	----------	-----------	---------------

	fsm_id [PK] bigint	stock_id integer	minute_id bigint	open numeric (10,2)	high numeric (10,2)	low numeric (10,2)	volume integer
1	600004202311200931	600004	202311200931	0.00	10.92	10.89	647
2	600004202311200932	600004	202311200932	0.00	10.90	10.89	407
3	600004202311200933	600004	202311200933	0.00	10.91	10.86	2568
4	600004202311200934	600004	202311200934	0.00	10.91	10.90	310
5	600004202311200935	600004	202311200935	0.00	10.95	10.91	1727
6	600004202311200936	600004	202311200936	0.00	10.94	10.92	202
7	600004202311200937	600004	202311200937	0.00	10.94	10.93	1230
8	600004202311200938	600004	202311200938	0.00	10.93	10.93	217
9	600004202311200939	600004	202311200939	0.00	10.94	10.93	86

Time Cost for PostgreSQL:

#	Node	Timings		F
		Exclusive	Inclusive	
1.	→ Bitmap Heap Scan on Trading.fact_stock... Recheck Cond: ((fact_stock_minute.stock_id = Heap Blocks: exact=2	0.489 ms	7.586 ms	F
2.	→ Bitmap AND (cost=644.38..644.38 r...	0.004 ms	7.098 ms	F
3.	→ Bitmap Index Scan using idx_fs... Index Cond: (fact_stock_minute.stoc	2.269 ms	2.269 ms	F
4.	→ Bitmap Index Scan using idx_fs... Index Cond: ((fact_stock_minute.min	4.825 ms	4.825 ms	F

Spark:

```

import time
start_time= time.time()
fact_stock_minute_df = spark.read.csv('Fact_Stock_Minute.csv', header=True, inferSchema=True)
stock_600004_info_df = fact_stock_minute_df.filter((col("Minute_ID") > 202311200930) \
                                                    & (col("Minute_ID") < 202311200940) \
                                                    & (col("Stock_ID") == 600004))

stock_600004_info_df.show()
print(f"Time used: {time.time() - start_time}")

```

✓ 4.8s Python

[Stage 29:> (0 + 3) / 3]

FSM_ID	Stock_ID	Minute_ID	Open	High	Low	Volume	Stock_Name	Turnover	Close
600004202311200931	600004	202311200931	0.0	10.92	10.89	647	0002000	705561.0	10.91
600004202311200932	600004	202311200932	0.0	10.9	10.89	401	0002000	436909.0	10.89
600004202311200933	600004	202311200933	0.0	10.91	10.86	2568	0002000	2794153.0	10.9
600004202311200934	600004	202311200934	0.0	10.91	10.9	310	0002000	338083.0	10.9
600004202311200935	600004	202311200935	0.0	10.95	10.91	1727	0002000	1886825.0	10.94
600004202311200936	600004	202311200936	0.0	10.94	10.92	202	0002000	220870.0	10.93
600004202311200937	600004	202311200937	0.0	10.94	10.93	1230	0002000	1345150.0	10.93
600004202311200938	600004	202311200938	0.0	10.93	10.93	211	0002000	230657.0	10.93
600004202311200939	600004	202311200939	0.0	10.94	10.93	86	0002000	94037.0	10.94

Time used: 4.820951700210571

Seq4: show total minutes in 2024-03-11

PostgreSQL:

```

23 --show total minutes in 2024-03-11
24 SELECT count(Minute_ID) AS Total_Minutes
25 from "Trading".Dim_Minute
26 where DATE_TRUNC('day',CAST(Date AS timestamp))='2024-03-11';
27

```

Data Output Messages Explain X Notifications

	total_minutes
1	241

Spark:

```

dim_minute_df = spark.read.csv('Dim_Minute_2024.csv', header=True, inferSchema=True)
total_minutes_df = dim_minute_df.filter(col("Date").cast("date") == "2024-03-11")\
    .agg({"Minute_ID": "count"})\
    .withColumnRenamed("count(Minute_ID)", "Total_Minutes")

total_minutes_df.show()

```

✓ 0.9s Python

Total_Minutes
241

Seq5: show total trading days in October

PostgreSQL:

```

28 --show total trading days in February 2024
29 SELECT count(Date_ID) AS Total_Trading_Days
30 from "Trading".Dim_Date
31 where year = 2024 and month = 2;
32

```

Data Output Messages Explain X Notifications

total_trading_days
16

Spark:

```
dim_date_df = spark.read.csv('Dim_Date_2024.csv', header=True, inferSchema=True)
total_trading_days_df = dim_date_df.filter((col("year") == 2024) & (col("month") == 2))\
    .agg({"Date_ID": "count"})\
    .withColumnRenamed("count(Date_ID)", "Total_Trading_Days")
total_trading_days_df.show()
```

✓ 0.4s Python

Total_Trading_Days
16

Seq6: recent company events from September

PostgreSQL:

```
--show recent company events from September
SELECT de.Stock_ID, de.Event_Date, de.Event_Type, de.Current_Flag,
       de.Effective_Timestamp, de.Expire_Timestamp
FROM "Trading".Fact_Company_Event fce
JOIN "Trading".Dim_Event de
ON fce.Event_Dim_ID = de.Event_Dim_ID
where fce.Date_ID>20230901 and fce.Stock_ID>600000 and fce.Stock_ID<600010;
```

--select Minute Level Data from 2024-03-11 9:30 to 2024-03-11 9:40
where the company is in CN market and its Total_Value is in top 10

Data Output Messages Explain × Notifications

stock_id	event_date	event_type	current_flag	effective_timestamp	expire_timestamp
integer	date	character varying (50)	boolean	timestamp without time zone	timestamp without time zone
600006	2023-09-15	Asset Restructuring	true	2023-09-15 00:00:00	[null]
600008	2023-11-14	Asset Restructuring	true	2023-11-14 00:00:00	[null]

Spark:

```
fact_company_event_df = spark.read.csv('Fact_Company_Event.csv', header=True, inferSchema=True).withColumnRenamed("Stock_ID", "FCE_Stock_ID")
dim_event_df = spark.read.csv('Dim_Event.csv', header=True, inferSchema=True).withColumnRenamed("Stock_ID", "DE_Stock_ID")

recent_events_df = fact_company_event_df.join(dim_event_df, fact_company_event_df.Event_Dim_ID == dim_event_df.Event_Dim_ID, "inner")
recent_events_df = recent_events_df.select("FCE_ID", "FCE_Stock_ID", "Date_ID", "Event_Type", "Current_Flag", "Effective_Timestamp", "Expire_Timestamp")
recent_events_df = recent_events_df.filter((col("Date_ID") > 20230901) & (col("FCE_Stock_ID") > 600000) & (col("FCE_Stock_ID") < 600010))
recent_events_df.show()
```

✓ 0.4s Python

FCE_ID	FCE_Stock_ID	Date_ID	Event_Type	Current_Flag	Effective_Timestamp	Expire_Timestamp
202309150002600006	600006	20230915	Asset Restructuring	1	2023-09-15	NULL
202311140063600008	600008	20231114	Asset Restructuring	1	2023-11-14	NULL

Seq7 select Minute Level Data from 2023-11-21 9:30 to 2023-11-21 9:40

where the company is in CN market and its Total_Value is in top 10

PostgreSQL:

```
43 SELECT fsm.FSM_ID, fsm.Stock_ID, fsm.Minute_ID, fsm.Close, fsm.High, fsm.Low, fsm.volu
44 FROM "Trading".Fact_Stock_Minute fsm
45 WHERE fsm.Stock_ID in (
46     SELECT Stock_ID
47     FROM "Trading".Dim_Stock
48     WHERE Stock_Market = 'CN'
49     ORDER BY Total_Value DESC
50     LIMIT 10
51 )
52 AND fsm.Minute_ID >= 202311210930 AND fsm.Minute_ID <= 202311210940;
53
54
```

Data Output Messages Explain X Notifications							
	fsm_id [PK] bigint	stock_id integer	minute_id bigint	close numeric (10,2)	high numeric (10,2)	low numeric (10,2)	volume integer
1	600938202311210930	600938	202311210930	19.12	19.12	19.12	
2	600938202311210931	600938	202311210931	19.06	19.14	19.05	4
3	600938202311210932	600938	202311210932	19.05	19.09	19.04	3
4	600938202311210933	600938	202311210933	19.01	19.05	19.00	3
5	600938202311210934	600938	202311210934	19.04	19.04	18.99	3
6	600938202311210935	600938	202311210935	19.05	19.05	19.04	3
7	600938202311210936	600938	202311210936	19.01	19.04	19.01	2
8	600938202311210937	600938	202311210937	19.02	19.02	19.01	2
9	600938202311210938	600938	202311210938	19.02	19.03	19.01	1
10	600938202311210939	600938	202311210939	19.02	19.03	19.01	1

Time Cost: 119.978ms

#	Node	Timings		Rows		
		Exclusive	Inclusive	Rows X	Actual	Plan
1.	→ Nested Loop Inner Join (cost=1371.75...	95.093 ms	119.978 ms	↓ 1.05	110	105
2.	→ Aggregate (cost=581.38..581.48 r... Buckets: Batches: Memory Usage: 24 k	0.157 ms	14.365 ms	↑ 1	10	10
3.	→ Limit (cost=581.23..581.25 r...	0.001 ms	14.209 ms	↑ 1	10	10
4.	→ Sort (cost=581.23..594....	1.151 ms	14.208 ms	↑ 530.9	10	5309
5.	→ Seq Scan on Tradin... Filter: ((dim_stock.stoc Rows Removed by Filte	13.057 ms	13.057 ms	↑ 1	5309	5309

Spark:


```

start_time= time.time()
top_10_cn_stocks_cached = dim_stock_df.filter(col("Stock_Market") == "CN") \
.orderBy(col("Total_Value").desc()) \
.limit(10) \
.select("Stock_ID") \
.withColumnRenamed("Stock_ID","top10_Stock_ID")\
.cache()

minute_level_data_df = fact_stock_minute_df.join(top_10_cn_stocks_cached, \
fact_stock_minute_df.Stock_ID == top_10_cn_stocks_cached.top10_Stock_ID, "inner") \
.filter((col("Minute_ID") >= 202311210930) & (col("Minute_ID") <= 202311210940)) \
.select("FSM_ID", "Stock_ID", "Minute_ID", "Close", "High", "Low", "volume")

minute_level_data_df.show()

top_10_cn_stocks_cached.unpersist()
print(f"Time used: {time.time() - start_time}")

```

✓ 0.6s Python

FSM_ID	Stock_ID	Minute_ID	Close	High	Low	volume
601288202311210930	601288	202311210930	3.66	3.66	3.66	11861
601288202311210931	601288	202311210931	3.66	3.67	3.65	155799
601288202311210932	601288	202311210932	3.66	3.67	3.65	116492
601288202311210933	601288	202311210933	3.67	3.67	3.66	11392
601288202311210934	601288	202311210934	3.66	3.67	3.66	14070
601288202311210935	601288	202311210935	3.67	3.67	3.66	28420
601288202311210936	601288	202311210936	3.66	3.67	3.66	14610
601288202311210937	601288	202311210937	3.67	3.67	3.66	9740
601288202311210938	601288	202311210938	3.66	3.67	3.66	9401
601288202311210939	601288	202311210939	3.66	3.67	3.66	10829
601288202311210940	601288	202311210940	3.67	3.67	3.66	38237
601857202311210930	601857	202311210930	7.14	7.14	7.14	4878
601857202311210931	601857	202311210931	7.15	7.15	7.12	23906
601857202311210932	601857	202311210932	7.12	7.14	7.12	9806
601857202311210933	601857	202311210933	7.12	7.13	7.11	4412
601857202311210934	601857	202311210934	7.14	7.14	7.12	5834
601857202311210935	601857	202311210935	7.15	7.15	7.13	14307
601857202311210936	601857	202311210936	7.14	7.15	7.14	7101
601857202311210937	601857	202311210937	7.14	7.15	7.14	2217
601857202311210938	601857	202311210938	7.15	7.15	7.14	4517

only showing top 20 rows

Time used: 0.6944880485534668

Time Cost: 0.69ms

PostgreSQL Performance Tuning Example1

select Daily Situations for every Monday in November of company SZ300796;

version 1 SELECT dd.Month, dd.WeekDay, fsd.Stock_ID, fsd.Date_ID,

fsd.Open,fsd.High,fsd.Low, fsd.Volume, fsd.Turnover

FROM "Trading".Fact_Stock_Daily fsd

JOIN "Trading".Dim_Date dd ON fsd.Date_ID = dd.Date_ID

WHERE dd.Month = 11 and dd.WeekDay=1 and fsd.Stock_ID = 300796;

Data Output Messages Explain X Notifications									
Graphical Analysis Statistics				Timings		Rows			
#	Node	Exclusive	Inclusive	Rows X	Actual	Plan	Loops		
1.	→ Gather (cost=1005.68.22327.38 rows=4 width=46) (actual=29.196.229.214 rows=4 loop...	6.384 ms	229.214 ms	↑ 1	4	4	1		
2.	→ Hash Inner Join (cost=5.68.21326.98 rows=2 width=46) (actual=142.25.222.83 row... Hash Cond: (fsd.date_id = dd.date_id)	148.608 ms	222.83 ms	↑ 2	1	2	3		
3.	→ Seq Scan on Trading_fact_stock_daily as fsd (cost=0.21321.07 rows=90 width=... Filter: (fsd.stock_id = 300796) Rows Removed by Filter: 352637	74.172 ms	74.172 ms	↑ 0.41	74	90	3		
4.	→ Hash (cost=5.63.5.63 rows=4 width=12) (actual=0.05.0.051 rows=4 loops=2) Buckets: 1024 Batches: 1 Memory Usage: 9 kB	0.032 ms	0.051 ms	↑ 0.34	4	4	2		
5.	→ Seq Scan on Trading_dim_date as dd (cost=0.5.63 rows=4 width=12) (actu... Filter: ((dd.month = 11) AND (dd.weekday = 1)) Rows Removed by Filter: 238	0.019 ms	0.019 ms	↑ 0.5	4	4	2		

version 2 SELECT dd.Month, dd.WeekDay, fsd.Stock_ID, fsd.Date_ID, fsd.Open,

fsd.High,fsd.Low, fsd.Volume, fsd.Turnover

FROM "Trading".Fact_Stock_Daily fsd

JOIN (SELECT Date_ID, Month, WeekDay FROM "Trading".Dim_Date

WHERE Month = 11 AND WeekDay = 1) dd ON fsd.Date_ID = dd.Date_ID

WHERE fsd.Stock_ID = 300796;

Data Output Messages Explain X Notifications									
Graphical Analysis Statistics				Timings		Rows			
#	Node	Exclusive	Inclusive	Rows X	Actual	Plan	Loops		
1.	→ Gather (cost=1005.68.22331.18 rows=4 width=46) (actual=228.062.229.117 rows=4 loo...	23.648 ms	229.117 ms	↑ 1	4	4	1		
2.	→ Hash Inner Join (cost=5.68.21330.78 rows=2 width=46) (actual=148.778.205.469 r... Hash Cond: (fsd.date_id = dim.date_id)	137.016 ms	205.469 ms	↑ 2	1	2	3		
3.	→ Seq Scan on Trading_fact_stock_daily as fsd (cost=0.21324.86 rows=90 width=... Filter: (fsd.stock_id = 300796) Rows Removed by Filter: 382637	68.4 ms	68.4 ms	↑ 0.41	74	90	3		
4.	→ Hash (cost=5.63.5.63 rows=4 width=12) (actual=0.052.0.053 rows=4 loops=3) Buckets: 1024 Batches: 1 Memory Usage: 9 kB	0.038 ms	0.053 ms	↑ 0.34	4	4	3		
5.	→ Seq Scan on Trading_dim_date as dim_date (cost=0.5.63 rows=4 width=12... Filter: ((dim_date.month = 11) AND (dim_date.weekday = 1)) Rows Removed by Filter: 238	0.015 ms	0.015 ms	↑ 0.34	4	4	3		

version 3 Add index on FK

CREATE INDEX idx_fsd_stock_id ON "Trading".Fact_Stock_Daily(Stock_ID);

CREATE INDEX idx_fsd_date_id ON "Trading".Fact_Stock_Daily(Date_ID);

Data Output Messages Explain X Notifications									
Graphical Analysis Statistics				Timings		Rows			
#	Node	Exclusive	Inclusive	Rows X	Actual	Plan	Loops		
1.	→ Nested Loop Inner Join (cost=66.47..287.61 rows=4 width=46) (actual=0.373..1.263 rows=4 L...	0.938 ms	1.263 ms	↑ 1	4	4	1		
2.	→ Seq Scan on Trading.dim_date as dim_date (cost=0..5.63 rows=4 width=12) (actual=0.01...	0.015 ms	0.015 ms	↑ 1	4	4	1		
3.	→ Bitmap Heap Scan on Trading.fact_stock_daily as fsd (cost=66.47..70.48 rows=1 width=3...	0.003 ms	0.31 ms	↑ 1	1	1	4		
4.	→ Bitmap AND (cost=66.47..66.47 rows=1 width=0) (actual=0.308..0.308 rows=0 loops...	0.23 ms	0.308 ms	↓ 1	0	1	4		
5.	→ Bitmap Index Scan using idx_fsd_stock_id (cost=0..6.05 rows=217 width=0) (ac...	0.002 ms	0.002 ms	↓ 0.26	222	217	4		
6.	→ Bitmap Index Scan using idx_fsd_date_id (cost=0..60.11 rows=5291 width=0) (a...	0.076 ms	0.076 ms	↓ 0.26	5302	5291	4		

Performance tuning skills:

Use subquery and Add index on FKs

PostgreSQL Performance Tuning Example2

Query: In event, we first filter out the latest events from 2023-11-20 to 2023-11-24, and from these events, we check the stock price changes of the corresponding companies on the next day, and from the companies that have increased by more than 3%, we categorize the companies according to the time and the type of the event, and finally, we sort them according to the date, and the average growth of the stock price.

Version 1:

with event_stock as(

select distinct ds.stock_id, fce.Date_ID, de.Event_Type

from "Trading".Fact_Company_Event fce

join "Trading".dim_event de on fce.Event_Dim_ID = de.Event_Dim_ID

join "Trading".Dim_Stock ds on ds.Stock_ID = fce.Stock_ID

where de.Expire_Timestamp isnull

and fce.Date_ID between 20231120 and 20231124

order by fce.Date_ID)

select distinct tmp.event_type, tmp.date_id, tmp.avg_incre

from (select

fsd.stock_id, fsd.date_id, es.event_type,fsd.percent_of_incre_decre,

AVG(fsd.percent_of_incre_decre) OVER(PARTITION BY es.event_type) avg_incre

from "Trading".Fact_Stock_Daily as fsd

join event_stock es on fsd.Stock_ID = es.stock_id and

fsd.Date_ID = es.Date_ID + 1 and fsd.percent_of_incre_decre>3.0

group by fsd.stock_id,fsd.date_id, es.event_type,fsd.percent_of_incre_decre

) AS tmp

group by date_id,event_type,tmp.avg_incre

order by tmp.date_id,tmp.avg_incre

#	Node	Exclusive	Inclusive
1.	→ Unique (cost=5661.31..5661.45 rows=8 width=54) (actual=227.707..227.8...	0.003 ms	227.889 ms
2.	→ Group (cost=5661.31..5661.39 rows=8 width=54) (actual=227.706..2...	0.004 ms	227.887 ms
3.	→ Sort (cost=5661.31..5661.33 rows=8 width=54) (actual=227.706...	0.01 ms	227.884 ms
4.	→ Subquery Scan (cost=5660.97..5661.19 rows=8 width=54) (...)	0.002 ms	227.874 ms
5.	→ Window Aggregate (cost=5660.97..5661.11 rows=8 wi...	0.176 ms	227.873 ms
6.	→ Sort (cost=5660.97..5660.99 rows=8 width=32) (a...	0.581 ms	227.697 ms
7.	→ Group (cost=5660.75..5660.85 rows=8 width=...	0.007 ms	227.116 ms

Version2

```
CREATE INDEX IF NOT EXISTS idx_fsm_stock_id ON "Trading".Fact_Stock_Minute(Stock_ID);
```

```
CREATE INDEX IF NOT EXISTS idx_fsm_minute_id ON "Trading".Fact_Stock_Minute(Minute_ID);
```

```
CREATE INDEX IF NOT EXISTS idx_fce_stock_id ON "Trading".Fact_Company_Event(Stock_ID);
```

```
CREATE INDEX IF NOT EXISTS idx_fce_date_id ON "Trading".Fact_Company_Event(Date_ID);
```

```
WITH event_stock AS (
```

```
    SELECT distinct ds.stock_id, fce.Date_ID, de.Event_Type
```

```
    FROM "Trading".Fact_Company_Event fce
```

```
    JOIN "Trading".dim_event de ON fce.Event_Dim_ID = de.Event_Dim_ID
```

```
    JOIN "Trading".Dim_Stock ds ON ds.Stock_ID = fce.Stock_ID
```

```
    WHERE de.Expire_Timestamp isnull
```

```
    AND fce.Date_ID between 20231120 and 20231124
```

```
),
```

```
tmp AS (
```

```
    SELECT
```

```
        fsd.stock_id, fsd.date_id, es.event_type,
```

```
        AVG(fsd.percent_of_incre_decre) OVER(PARTITION BY es.event_type) avg_incre
```

```
    FROM "Trading".Fact_Stock_Daily as fsd
```

JOIN event_stock es ON

fsd.Stock_ID = es.stock_id and

fsd.Date_ID = es.Date_ID + 1 and

fsd.percent_of_incre_decre>3.0

GROUP BY fsd.stock_id,fsd.date_id, es.event_type,fsd.percent_of_incre_decre

)

SELECT distinct

tmp.event_type,

tmp.date_id,

tmp.avg_incre

FROM tmp

GROUP BY date_id,event_type,tmp.avg_incre

ORDER BY tmp.date_id,tmp.avg_incre

		Exclusive	Inclusive	R
1.	→ Unique (cost=5662.26..5662.4 rows=8 width=54) (actual=62.855..62.889 r...	0.003 ms	62.889 ms	
2.	→ Group (cost=5662.26..5662.34 rows=8 width=54) (actual=62.855..62....	0.004 ms	62.887 ms	
3.	→ Sort (cost=5662.26..5662.28 rows=8 width=54) (actual=62.854.....	0.016 ms	62.883 ms	,
4.	→ Subquery Scan (cost=5661.92..5662.14 rows=8 width=54) (...)	0.004 ms	62.868 ms	,
5.	→ Window Aggregate (cost=5661.92..5662.06 rows=8 wi...	0.304 ms	62.865 ms	,
6.	→ Sort (cost=5661.92..5661.94 rows=8 width=32) (a...	0.525 ms	62.562 ms	,
7.	→ Group (cost=1686.93..5661.8 rows=8 width=3...	0.006 ms	62.037 ms	,
8.	→ Incremental Sort (cost=1686.93..5661.7...	0.029 ms	62.032 ms	,

Explanation:

First of all, add index to all foreign keys, and put all fk's into block buffer cache to improve cache hit ratio. meanwhile, adding the attribute of index can avoid full table scan when joining.

Use the WITH clause to create a temporary table instead of duplicating the event_stock and Fact_Stock_Daily tables in the SELECT clause. This avoids double counting and improves query speed and makes the size of the index scan during a join operation significantly smaller

Instead of using a global sort, I use the PARTITION BY clause to group data by date or other dimension. This reduces the overhead of sorting operations and improves query efficiency.

Instead of using multiple subqueries, I use the AVG() function to calculate the average increment for each date. This simplifies the query logic and improves query performance.